

Data Structures Final Project Analysis


Tyler Andrews
December 16, 2016

All in all, this assignment was very eye opening. I was extremely surprised by the differences in performance between the sorting algorithms we tried, and specifically how much more efficient quicksort was than the others. On average, it was a few hundred times faster for large inputs. It managed to sort 50,000 numbers in under 10ms, which was much faster than I would have ever guessed. Here you can see my program's output for an input file containing 50,000 randomized doubles:

Quick sort initiated at time: 10000	Insertion sort initiated at time: 20000
Quick sort finished at time: 20000	Insertion sort finished at time: 2980000
Quick Sort Complete. Run Time: 10ms	Insertion Sort Complete. Run Time: 2960ms

In this case, quicksort was nearly 300 times faster than insertion sort, and (not shown) 600 times faster than gnome sort. This is a good representation of logarithmic vs. quadratic performance. On paper, $O(n^2)$ may not look much larger than $O(n \log n)$, but this exercise certainly shows how far apart they really are in terms of speed and efficiency.

As far as performance, the brute force algorithms had a very clear hindrance on my VM. Both insertion sort and gnome sort spiked CPU usage for my VM to 100%. Here you can see the absurd spike in resources when insertion sort is initiated. The program also appears to use about 800MB of RAM, which seems very high.



I think this analysis would have benefited from a visual component other than seeing the run times themselves. Seeing the data being sorted would be a great way to show the speed of the algorithms (like the YouTube videos we watched in class) in a more direct way. Overall, I was extremely impressed by the efficiency of quicksort, and I think this assignment was a very good indicator of the value of recursion.